

Analysis of AirBnB data

I used SQL functions for my assignment – although I did familiarize myself with map/reduce/filter etc. in the process. My task was then mostly transforming the data into a format I could use the correct SQL functions on. I have not put a lot of emphasis on explaining why I use SQL functions like AVG and SUM, as I think most are quite familiar with them. I split each task into a separate .scala file - the solution could be more efficient if the DataFrames were reused between the tasks, but I thought being able to run each task individually was more important.

About task 6:

I managed to find a library function “pointInPolygon” and get my DataFrame on the form

```
[Neighbourhood name, [WrappedArray[WrappedArray[WrappedArray[....]]]]
```

I spent several hours trying to figure out how to get the coordinates out of the wrapped array, but eventually ran out of time. I have provided the code for as far as I got. My procedure would be for each listing, check if the latitude and longitude is within a polygon and then use the neighborhood name.

I also had issues exporting outputs to csv and therefore used println and copied this to a text file.

2a

```
CREATE TABLE listings
```

```
...
```

```
id int NOT NULL PRIMARY KEY
```

```
...
```

```
CREATE TABLE reviews
```

```
...
```

```
FOREIGN KEY (listing_id) REFERENCES listings(id)
```

```
...
```

```
CREATE TABLE reviews
```

```
...
```

```
FOREIGN KEY (listing_id) REFERENCES listings(id)
```

```
...
```

2b

I used `spark.sql("select count(distinct " + column + ") from listings")` and looped through the columns. The output for this task is available in the `outputs.txt` file.

2c

3 cities: New York, Seattle and San Francisco

I used `spark.sql("select distinct city from listings")`

2d

`id`: primary key used for further analysis by joining with other tables as foreign key

`price`: booking price per night in USD, min \$0 max none

`number_of_reviews`: self-explanatory, min 0, max none

`latitude`, `longitude`: location information

Task 3

I transformed the original DataFrame where price was a string to a new DataFrame with price as an int so I could use SQL functions like avg and sum using a user defined function (udf):

```
spark.udf.register("myConvertCurrency", (input: String) =>  
java.text.NumberFormat.getCurrencyInstance(java.util.Locale.US).parse(input).intValue.toInt)
```

with

```
listingsRaw.withColumn("price", callUDF("myConvertCurrency", listingsRaw("price")))
```

This transformation was also used to solve all subsequent tasks.

For each city, I created a view "citySubset" for better reuse as such:

```
spark.sql("select price, reviews_per_month, room_type from listings where city = '" + city + "'")
```

3a

City	Average price
San Francisco	\$251
Seattle	\$131
New York	\$149

I used `spark.sql("select avg(price) from citySubset")`

3b

City	Shared room	Entire home/apt	Private room
San Francisco	\$88	\$341	\$135
Seattle	\$52	\$159	\$77
New York	\$69	\$208	\$90

I obtained the distinct room_types as such:

```
spark.sql("select distinct room_type from listings")
```

Then I used

```
spark.sql("select avg(price) from citySubset where room_type = '" + roomType + "'")
```

3c

City	Average number of reviews (per listing)
San Francisco	1.68
Seattle	2.08
New York	1.38

I used

```
spark.sql("select avg(reviews_per_month) from citySubset")
```

3d

City	Estimated number of booked nights per year
San Francisco	196512
Seattle	113723
New York	923823

I defined a udf

```
spark.udf.register("myNightsPerYear", (input: String) => ((if(input == null) 0.0 else input.toDouble) / 0.7) * 12)
```

which I applied to my "citySubset" to obtain a new column "nights_per_year"

```
citySubset.withColumn("nights_per_year", callUDF("myNightsPerYear",  
citySubset("reviews_per_month"))).createOrReplaceTempView("citySubsetNightsPerYear")
```

I then used

```
spark.sql("select sum(nights_per_year) from citySubsetNightsPerYear")
```

3e

City	Estimated amount of money spent per year (\$)
San Francisco	$3.38 \cdot 10^7$
Seattle	$1.25 \cdot 10^7$
New York	$1.26 \cdot 10^8$

I reused "citySubsetNightsPerYear" from the previous task and used

```
spark.sql("select sum(nights_per_year * price) from citySubsetNightsPerYear")
```

Task 4

I read the csv file and first created a view "listingsPerHost", containing number of listings per host:

```
spark.sql("select count(*) as count from listings group by host_id")
```

4a

Average number of listings per host: 1.26

I used

```
spark.sql("select avg(count) from listingsPerHost")
```

4b

Percentage of hosts with more than 1 listing: 15%

I used

```
spark.sql("select count from listingsPerHost where count > 1").count.toFloat / listingsPerHost.count
```

since listingsPerHost.count is equal to the number of unique hosts

4c

City	Hosts (host_name, income)
San Francisco	Max (\$3950000) Ramil (\$3920000) Matt (\$1913800)
Seattle	Jordan (\$3420683) Sea To Sky Rentals (\$3330909) Daniela (\$2487022)
New York	Jessica & Doug (\$7421883) John (\$4811400) 124 (\$4303309)

I created a view "countsPerListings" by joining listings and calendar on the listings.id = calendar.listing_id where available was set to "t" (I assume this means booked), grouping by the listing_id which then represents amount of nights a listing was booked:

```
spark.sql("select count(*) as count, listing_id from calendar JOIN listings ON listings.id =  
calendar.listing_id where available = 't' and city = '" + city + "'" + "group by listing_id")
```

I then found the income for each host by grouping on host_id and using sum(price * countsPerListings.count), ordering by the income descending and limiting the output to 3:

```
spark.sql("select first(host_name), sum(price * countsPerListings.count) as income from listings join  
countsPerListings on listings.id = countsPerListings.listing_id group by host_id order by income desc  
limit 3")
```

Task 5

I think the task text could be a bit clearer here – I'm not quite sure what the relationship between a review and nights spent is so I'm going to assume that an entry in the reviews table represents 3 nights spent for the respective listing.

I first created a view "reviewsListings", joining reviews and listings tables and selecting appropriate columns.

5a

City	Guests (reviewer_name, nights_spent)
San Francisco	Emily (87) Zafar (81) Claire (75)
Seattle	Amanda (213) Kathryn (102) David (81)
New York	J. B. (79) Andy (41) Adrienne (33)

I used

```
spark.sql("select first(reviewer_name), (count(*) * 3) as count from reviewsListings where city = '' + city + '' group by reviewer_id order by count desc limit 3")
```

Grouping by reviewer_id gives number of listings belonging to a reviewer.

5b

Guest who spent the most money: Claire (\$57132)

I used

```
spark.sql("select reviewer_id, first(reviewer_name), sum(price) * 3 as spent from reviewsListings group by reviewer_id order by spent desc limit 1")
```